

CUDA Fortran Dynamic Parallelism (1)

```
attributes(global) subroutine strassen( A, B, C, M, N, K )  
real :: A(M,K), B(K,N), C(M,N)  
integer, value :: N, M, K  
...  
if (ntimes == 0) then  
  allocate(m1(1:m/2,1:k/2))  
  allocate(m2(1:k/2,1:n/2))  
  allocate(m3(1:m/2,1:k/2))  
  allocate(m4(1:k/2,1:n/2))  
  allocate(m5(1:m/2,1:k/2))  
  allocate(m6(1:k/2,1:n/2))  
  allocate(m7(1:m/2,1:k/2))  
  flags = cudaStreamNoBlocking  
  do i = 1, 7  
    istat = cudaStreamCreateWithFlags(istreams(i), flags)  
  end do  
end if  
...
```

Support for Fortran allocate and deallocate in device code

CUDA Fortran Dynamic Parallelism (2)

```
...  
! m1 = (A11 + A22) * (B11 + B22)  
call dgemm16t1<<<devblocks,devthreads,0,istreams(1)>>>(a(1,1), &  
    a(1+m/2,1+k/2), m, &  
    b(1,1), b(1+k/2,1+n/2), k, &  
    m1(1,1), newn, newn, 1.0d0)  
! m2 = (A21 + A22) * B11  
call dgemm16t2<<<devblocks,devthreads,0,istreams(2)>>>(a(1+m/2,1), &  
    a(1+m/2,1+k/2), m, &  
    b(1,1), k, &  
    m2(1,1), newn, newn)  
  
...  
  
! m7 = (A12 - A22) * (B21 + B22)  
call dgemm16t1<<<devblocks,devthreads,0,istreams(7)>>>(a(1,1+k/2), &  
    a(1+m/2,1+k/2), m, &  
    b(1+k/2,1), b(1+k/2,1+n/2), k, &  
    m7(1,1), newn, newn, -1.0d0)
```

istat = cuda

...

Support for kernel launch from device code

CUDA Fortran Dynamic Parallelism (3)

```
...  
! C11 = m1 + m4 - m5 + m7  
call add16x4<<<1,devthreads,0,istreams(1)>>>(m1,m4,m5,m7,m/2,c(1,1),m,n/2)  
  
! C12 = m3 + m5  
call add16<<<1,devthreads,0,istreams(2)>>>(m3,m/2,m5,m/2,c(1,1+n/2),m,n/2)  
  
! C21 = m2 + m4  
call add16<<<1,devthreads,0,istreams(3)>>>(m2,m/2,m4,m/2,c(1+m/2,1),m,n/2)  
  
! C22 = m1 + m3 - m2 + m6  
call add16x4<<<1,devthreads,0,istreams(4)>>>(m1,m3,m2,m6,m/2,c(1+m/2,1+n/2),m,n/2)  
  
...  
  
end subroutine strassen
```

Compile using -Mcuda=rdc,cc35

CUDA Fortran Separate Compilation

Compile separate modules independently

```
% pgf90 -c -O2 -Mcuda=rdc ddfun90.cuf ddm90.cuf
```

Object files can be put into a library

```
% ar rc ddfunc.a ddfun90.o ddm90.o
```

Use the modules in device code in typical Fortran fashion

```
% cat main.cuf
```

```
    program main
```

```
    use ddm90
```

```
    ...
```

Link using pgf90 and the rdc option

```
% pgf90 -O2 -Mcuda=rdc main.cuf ddfunc.a
```

Performance Optimization


- **Host-device data transfers**
 - Page-locked transfers
 - Asynchronous transfers
- **Device memory**
 - Coalescing
 - Shared memory
 - Textures
- **Execution Configuration**
 - Thread-level parallelism
 - Instruction-level parallelism

Host-Device Transfers

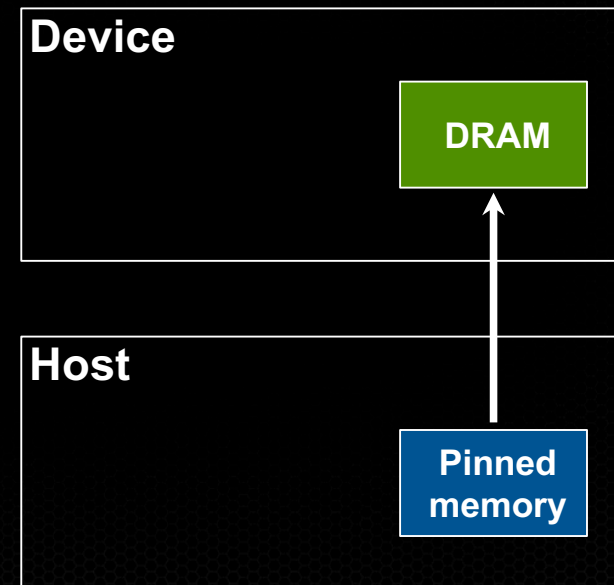
- Host-device bandwidth is much lower than bandwidth within device
 - 8 GB/s peak (PCIe x16 Gen 2) vs. 250 GB/s peak (Tesla K20X)
- Minimize number of transfers
 - Intermediate data can be allocated, used, and deallocated without copying to host memory
 - Sometimes better to do low parallelism operations on the GPU if it avoids transfers to and from host

Page-Locked Data Transfers

Pageable Data Transfer



Page-locked Data Transfer



Page-Locked Data Transfers

- Page-locked or pinned host memory by declaration
 - Designated by **pinned** variable attribute
 - Must be **allocatable**

```
real, device :: a_d(N)
real, pinned, allocatable :: a(:)
allocate(a(N), STAT=istat, PINNED=pinnedFlag)
...
a_d = a
```

- Tesla K20/Sandy Bridge
 - Pageable: ~3.3 GB/s
 - Pinned: ~6 GB/s

Overlapping Transfers and Computation

- Kernel launches are asynchronous, normal memory copies are blocking


```
a_d = a  ! blocks on host until transfer completes  
call inc<<<g,b>>>>(a_d, b) ! Control returns immediately to CPU  
a = a_d  ! starts only after kernel completes
```

- Asynchronous and Stream APIs allow overlap of transfers with computation
- A stream is a sequence of operations that execute in order on the GPU
 - Operations in different (*non-default*) streams can be interleaved
 - Stream ID used as arguments to async transfers and kernel launches

Asynchronous Data Transfers

- Asynchronous host-device transfers return control immediately to CPU
 - `cudaMemcpyAsync(dst, src, nElements, stream)`
 - Requires pinned host memory
- Overlapping data transfer with CPU computation
 - default stream = 0

```
istat = cudaMemcpyAsync(a_d, a_h, N, 0)
call kernel<<<grid, block>>>(a_d)
call cpuFunction(b)
```



Overlapping Transfers and Kernels

- Requires:

- Pinned host memory
- Kernel and transfer to use different *non-zero* streams

```
integer (kind=cuda_stream_kind) :: stream1, stream2
```

```
...
```

```
istat = cudaStreamCreate(stream1)
```

```
istat = cudaStreamCreate(stream2)
```

```
istat = cudaMemcpyAsync(a_d, a_h, N, stream1)
```

```
call kernel<<<grid, block, 0, stream2>>>(b_d)
```



overlapped

GPU/CPU Synchronization

- **cudaDeviceSynchronize()**
 - Blocks until all previously issued operations on the GPU complete
- **cudaStreamSynchronize(stream)**
 - Blocks until all previously issued operations to **stream** complete
- **cudaStreamQuery(stream)**
 - Indicates whether **stream** is idle
 - Does not block CPU code

GPU/CPU Synchronization

- Stream-based using CUDA events

- Events can be inserted into streams

```
type (cudaEvent) :: event
...
istat = cudaEventRecord(event, stream1)
```

- Event is recorded when the GPU reaches it in the stream
 - Recorded = assigned a time stamp
 - Useful for timing code
- **cudaEventSynchronize(event)**
 - Blocks CPU until event is recorded

Asynchronous Example

- `async.cuf`