

Performance Optimization

- **Host-device data transfers**
 - Page-locked transfers
 - Asynchronous transfers
- **Device memory**
 - Coalescing
 - Shared memory
 - Textures
- **Execution Configuration**
 - Thread-level parallelism
 - Instruction-level parallelism⁷

Coalescing

- Device data accessed by a group of 16 or 32 sequential threads can result in as little as a single transaction if certain requirements are met
 - Alignment
 - Stride
 - Generation of GPU architecture


Misaligned Data Access

- Use array increment kernel with variable misalignment

```
attributes(global) subroutine offset(a, s)
  real :: a(*)
  integer, value :: s
  integer :: i
  i = blockDim%x*(blockIdx%x-1)+threadIdx%x + s
  a(i) = a(i)+1
end subroutine offset
```

Array a() allocated with padding to avoid out-of-bounds accesses

Misaligned Data Access



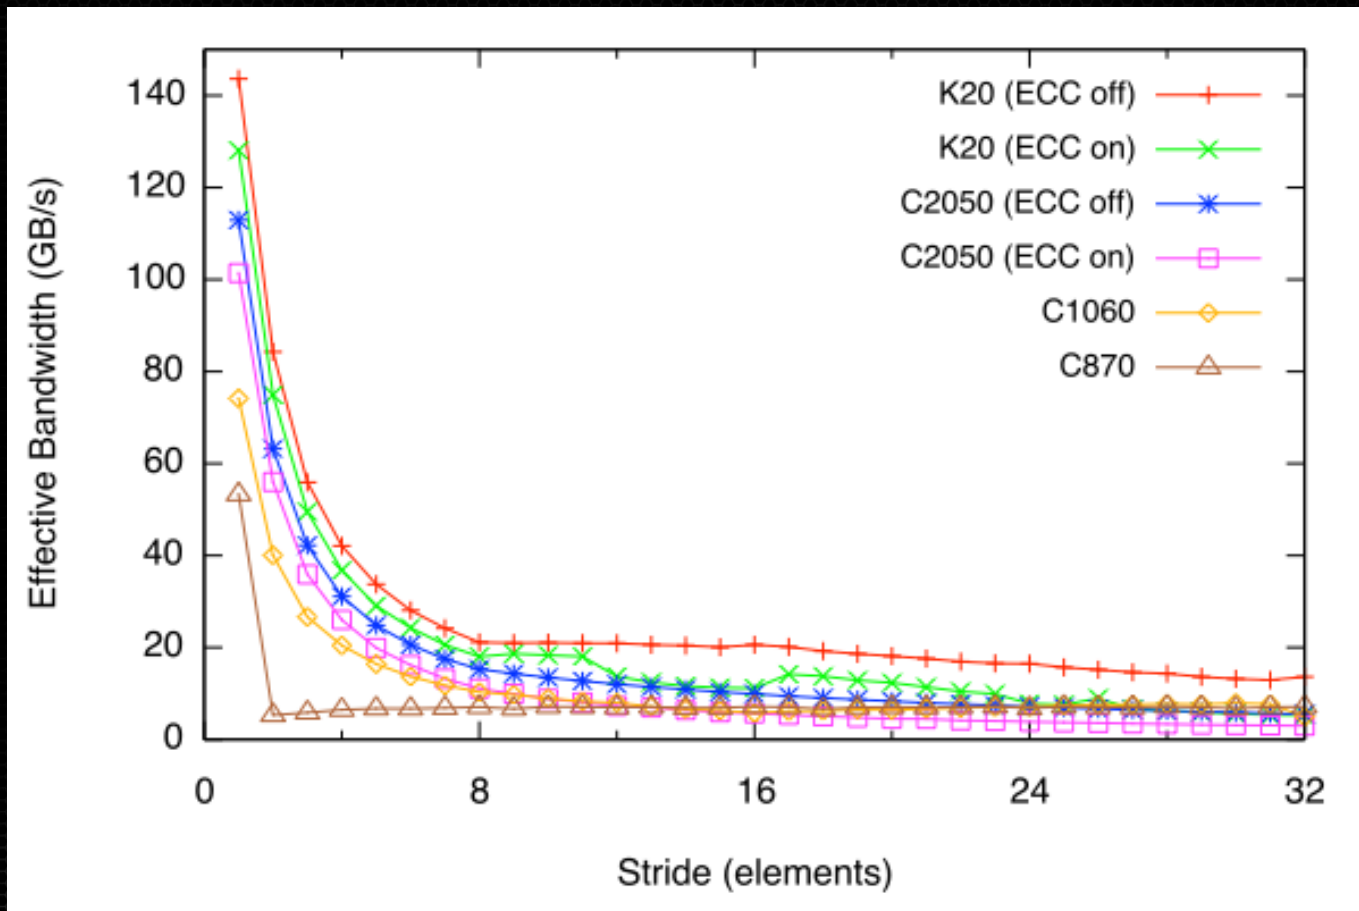
Strided Data Access

- Use array increment kernel with variable stride

```
attributes(global) subroutine stride(a, s)
  real :: a(*)
  integer, value :: s
  integer :: i
  i = (blockDim%x*(blockIdx%x-1)+threadIdx%x) * s
  a(i) = a(i)+1
end subroutine stride
```


Array a() allocated with padding to avoid out-of-bounds accesses

Strided Data Access



Shared Memory


- On-chip
- All threads in a block have access to same shared memory
- Used to reduce multiple loads of device data
- Used to accommodate coalescing



Matrix Transpose (transpose.cuf)

```
attributes(global) subroutine transposeNaive(odata, idata)
  real, intent(out) :: odata(ny,nx)
  real, intent(in)  :: idata(nx,ny)
  integer :: x, y

  x = (blockIdx%x-1) * blockDim%x + threadIdx%x
  y = (blockIdx%y-1) * blockDim%y + threadIdx%y
  odata(y,x) = idata(x,y)
end subroutine transposeNaive
```




Matrix Transpose - Shared Memory

```
attributes(global) subroutine transposeCoalesced(odata, idata)
  real, intent(out) :: odata(ny,nx)
  real, intent(in) :: idata(nx,ny)
  real, shared :: tile(TILE_DIM, TILE_DIM)
  integer :: x, y

  x = (blockIdx%x-1)*blockDim%x + threadIdx%x
  y = (blockIdx%y-1)*blockDim%y + threadIdx%y
  tile(threadIdx%x, threadIdx%y) = idata(x,y)

  call syncthreads()

  x = (blockIdx%y-1)*blockDim%y + threadIdx%x
  y = (blockIdx%x-1)*blockDim%x + threadIdx%y
  odata(x,y) = tile(threadIdx%y, threadIdx%x)
end subroutine transposeCoalesced
```



Shared Memory Bank Conflicts

- Shared memory is divided into banks that can be accessed simultaneously
- Successive 32-bit words are assigned to successive banks
- Requests for different data in the same bank by threads in a warp (SIMD unit of 32 successive threads) are serialized
- Common workaround is to pad shared memory arrays

```
real, shared :: tile(TILE_DIM +1, TILE_DIM)
```


Constant Memory

- Small (64KB), read-only in a kernel, written by the host
 - assignment or API
 - used for small coefficient tables, common pointers
- Hardware cached
 - separate 16KB constant cache

Textures

- Different pathway for accessing data in device DRAM
- Read-only by device code
- Cached on chip in texture cache
- Uses F90 pointer notation
 - Declare texture at module scope in module where used
`real, texture, pointer :: aTex(:)`
 - Bind and unbind texture in host code using pointer notation
- Equivalent to `tex1Dfetch()` in CUDA C
 - No filtering or wrapping modes

Texture Memory as a Read-only Cache

CUDA Fortran module

```
real, texture, pointer :: q(:)
```

CUDA Fortran host code

```
real, device, target :: p(:)
```

```
...
```

```
allocate(p(n))
```

```
q => p
```

CUDA Fortran device code

```
i = threadIdx%x
```

```
j = (blockIdx%x-1)*blockDim%x + i
```

```
s(j) = s(j) + q(i)
```

CUDA C/C++ global

```
texture<float, 1,  
    cudaReadModeElementType> tq;
```

CUDA C/C++ host code

```
float *d_p;
```

```
...
```

```
cudaMalloc((void **) d_p, size);
```

```
cudaChannelFormatDesc desc =
```

```
    cudaCreateChannelDesc(32, 0, 0,  
        0, cudaFormatKindFloat);
```

```
cudaBindTexture(0, tq, d_p,  
    desc, size);
```

CUDA C/C++ device code

```
i = threadIdx.x;
```

```
j = blockIdx.x*blockDim.x + i;
```

```
s[j] = s[j] + tex1Dfetch(tq, i);
```

Texture Memory Performance (Measured Gbytes/sec)

```

module memtests
integer, texture, pointer :: t(:)
contains
  attributes(global) subroutine stridem(m,b,s)
    integer, device :: m(*), b(*)
    integer, value :: s
    i = blockDim%x*(blockIdx%x-1) + threadIdx%x
    j = i * s
    b(i) = m(j) + 1
    return
  end subroutine

  attributes(global) subroutine stridet(b,s)
    integer, device :: b(*)
    integer, value :: s
    i = blockDim%x*(blockIdx%x-1) + threadIdx%x
    j = i * s
    b(i) = t(j) + 1
    return
  end subroutine
end module memtests

```

	blockDim = 32		blockDim = 128	
stride	stridem	stridet	stridem	stridet
1	53.45	50.94	141.93	135.43
2	49.77	49.93	100.37	99.61
3	46.96	48.07	75.25	74.91
4	44.25	45.90	60.19	60.00
5	39.06	39.89	50.09	49.96
6	37.14	39.33	42.95	42.93
7	32.88	35.94	37.56	37.50
8	32.48	32.98	33.38	33.42
9	29.90	32.94	30.01	33.38
10	27.43	32.86	27.28	33.07
11	25.17	32.77	24.98	32.91
12	23.23	33.19	23.07	33.01
13	21.57	33.13	21.40	32.26
14	20.15	32.98	19.97	31.76
15	18.87	32.80	18.72	30.80
16	17.78	32.66	17.60	31.83
20	14.38	33.37	14.23	26.84
24	12.08	33.71	11.94	24.30
28	10.41	33.38	10.27	19.97
32	9.15	33.42	9.02	20.19

Texture Example

- `strideTex.cuf`

Alternative to texture fetches

- Use INTENT(IN) as attribute for dummy arguments

```
attributes(global) subroutine test( a, b )  
integer, device :: a(*)  
integer, device, intent(in) :: b(*)
```

- Compile with -Mcuda=keepptx,cc35
 - Verify with search for ld.global.nc operations in ptx file
- Compile with -Mcuda=keepbin,cc35
 - Verify with “cuobjdump -dump-sass” on the .bin file

Performance Optimization

- **Host-device data transfers**
 - Page-locked transfers
 - Asynchronous transfers
- **Device memory**
 - Coalescing
 - Shared memory
 - Textures
- **Execution Configuration**
 - Thread-level parallelism
 - Instruction-level parallelism

Execution Configuration

- GPUs are high latency, 100s of cycles per device memory request
- For good performance, you need to ensure there is enough parallelism to hide this latency
- Such parallelism can come from:
 - Thread-level parallelism
 - Instruction-level parallelism

Thread-Level Parallelism

- Execution configuration dictates number of threads per block
 - Limit on number of threads per block for each architecture
- Number of concurrent blocks on a multiprocessor limited by
 - Register use per thread
 - Shared memory use per thread block
 - Limit on number of threads per multiprocessor
- Occupancy
 - Ratio of actual to maximum number of concurrent threads per multiprocessor

Thread-Level Parallelism

```
attributes(global) subroutine copy(odata, idata)
  real :: odata(*), idata(*), tmp
  integer :: i
  i = (blockIdx%x-1)*blockDim%x + threadIdx%x
  tmp = idata(i)
  odata(i) = tmp
end subroutine copy
```

Thread-Level Parallelism

- Run on K20
 - Maximum of 2048 concurrent threads per multiprocessor
 - Maximum of 16 concurrent blocks per multiprocessor
 - Maximum of 1024 threads per block

Thread Block Size	Occupancy	Bandwidth (GB/s)
32	0.25	96
64	0.5	125
128	1	136
256	1	137
512	1	137
1024	1	133

Thread-Level Parallelism

- Mimic high resource use

- Specify enough shared memory so only one thread block can reside on a multiprocessor at a time

```
call copy<<<grid, threadBlock, 0.9*smBytes>>>(b_d, a_d)
```

Thread Block	No Shared Memory		Shared Memory	
	Occupancy	Bandwidth	Occupancy	Bandwidth
32	0.25	96	0.016	8
64	0.5	125	0.031	15
128	1	136	0.063	29
256	1	137	0.125	53
512	1	137	0.25	91
1024	1	133	0.5	123

Instruction-Level Parallelism

- Have each thread process multiple elements

```
attributes(global) subroutine copy_ILP(odata, idata)
  real :: odata(*), idata(*), tmp(ILP)
  integer :: i,j

  i = (blockIdx%x-1)*blockDim%x*ILP + threadIdx%x
  do j = 1, ILP
    tmp(j) = idata(i+(j-1)*blockDim%x)
  enddo
  do j = 1, ILP
    odata(i+(j-1)*blockDim%x) = tmp(j)
  enddo
end subroutine copy_ILP
```


Instruction-Level Parallelism

Thread Block Size	No Shared Memory		Shared Memory		
	Occupancy	Bandwidth	Occupancy	Bandwidth No ILP	Bandwidth ILP = 4
32	0.25	96	0.016	8	26
64	0.5	125	0.031	15	50
128	1	136	0.063	29	90
256	1	137	0.125	53	125
512	1	137	0.25	91	140
1024	1	133	0.5	123	139

Calling CUBLAS from CUDA Fortran

- Module which defines interfaces to CUBLAS from CUDA Fortran
 - `use cublas`
- Interfaces in three forms
 - Overloaded BLAS interfaces that take device array arguments
 - call `saxpy(n, a_d, x_d, incx, y_d, incy)`
 - Legacy CUBLAS interfaces
 - call `cublasSaxpy(n, a_d, x_d, incx, y_d, incy)`
 - Multi-GPU version (CUDA 4.0) that utilizes a handle `h`
 - `istat = cublasSaxpy_v2(h, n, a_d, x_d, incx, y_d, incy)`
- Mixing the three forms is allowed

Calling CUBLAS from CUDA Fortran

```
program cublasTest
  use cublas
  implicit none

  real, allocatable :: a(:, :), b(:, :), c(:, :)
  real, device, allocatable :: a_d(:, :), b_d(:, :), c_d(:, :)
  integer :: k=4, m=4, n=4
  real :: alpha=1.0, beta=2.0, maxError

  allocate(a(m,k), b(k,n), c(m,n), a_d(m,k), b_d(k,n), c_d(m,n))

  a = 1; a_d = a
  b = 2; b_d = b
  c = 3; c_d = c

  call cublasSgemm('N', 'N', m, n, k, alpha, a_d, m, b_d, k, beta, c_d, m) ! or sgemm(...)

  c=c_d
  write(*,*) 'Maximum error: ', maxval(abs(c-14.0))

  deallocate (a,b,c,a_d,b_d,c_d)

end program cublasTest
```

CUDA Fortran Profiling and Debugging

- Some examples of profiling CUDA Fortran and the current status of CUDA Fortran debugging.

Performance Measurement

- CUDA Event management
- pgcollect / pgprof
- nvvp visual profiler; nvprof text profiler
- others (TAU, Vampir, ...)

Event Management

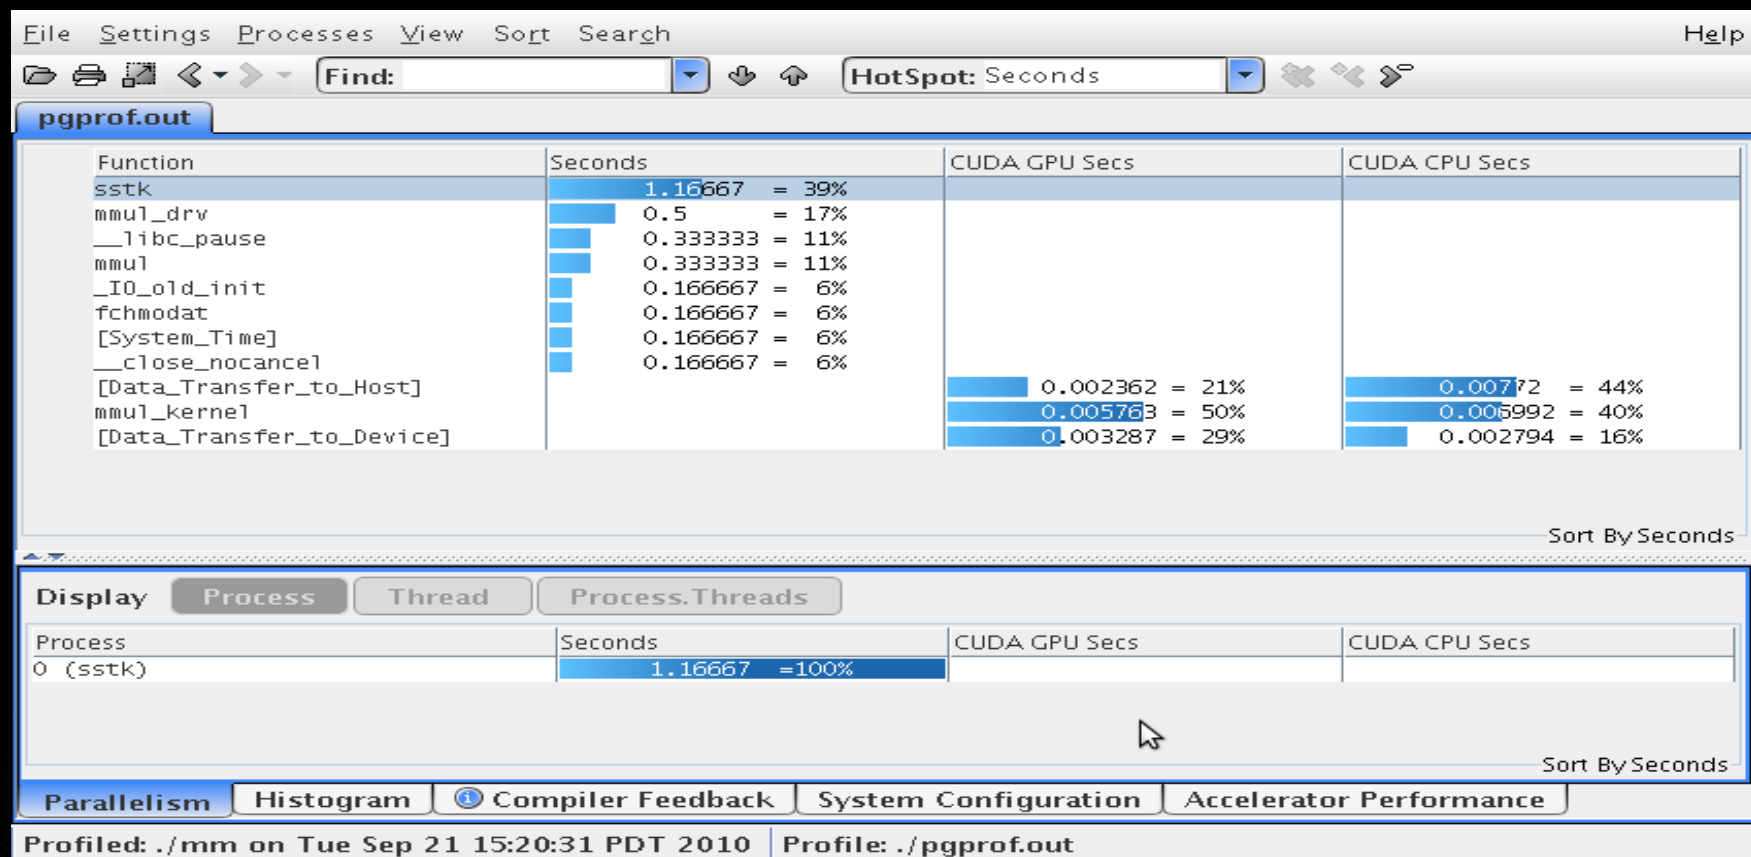
- `integer cudaEventCreate(event)`
`type(cudaEvent), intent(out) :: event`
- `integer cudaEventRecord(event, stream)`
`type(cudaEvent), intent(in) :: event`
`integer, intent(in) :: stream`
- `integer cudaEventSynchronize(event)`
`type(cudaEvent), intent(in) :: event`
- `integer cudaEventElapsedTime(time,e1,e2)`
`real, intent(out) :: time`
`type(cudaEvent), intent(in) :: e1, e2`
- `integer cudaEventDestroy(event)`
`type(cudaEvent), intent(in) :: event`

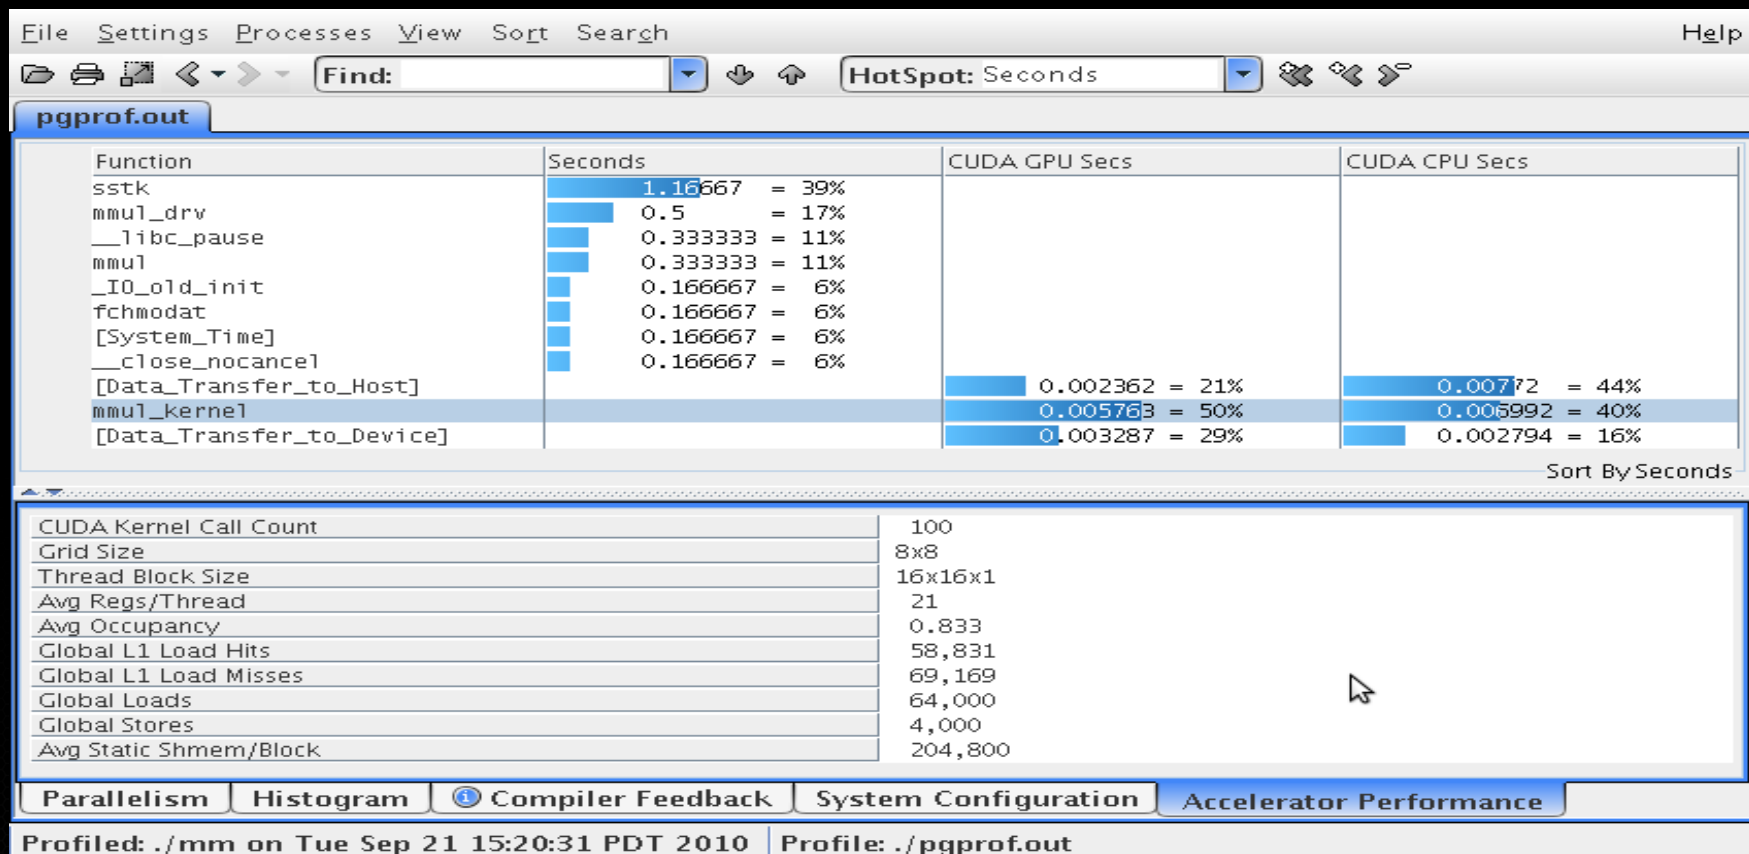
Events Example

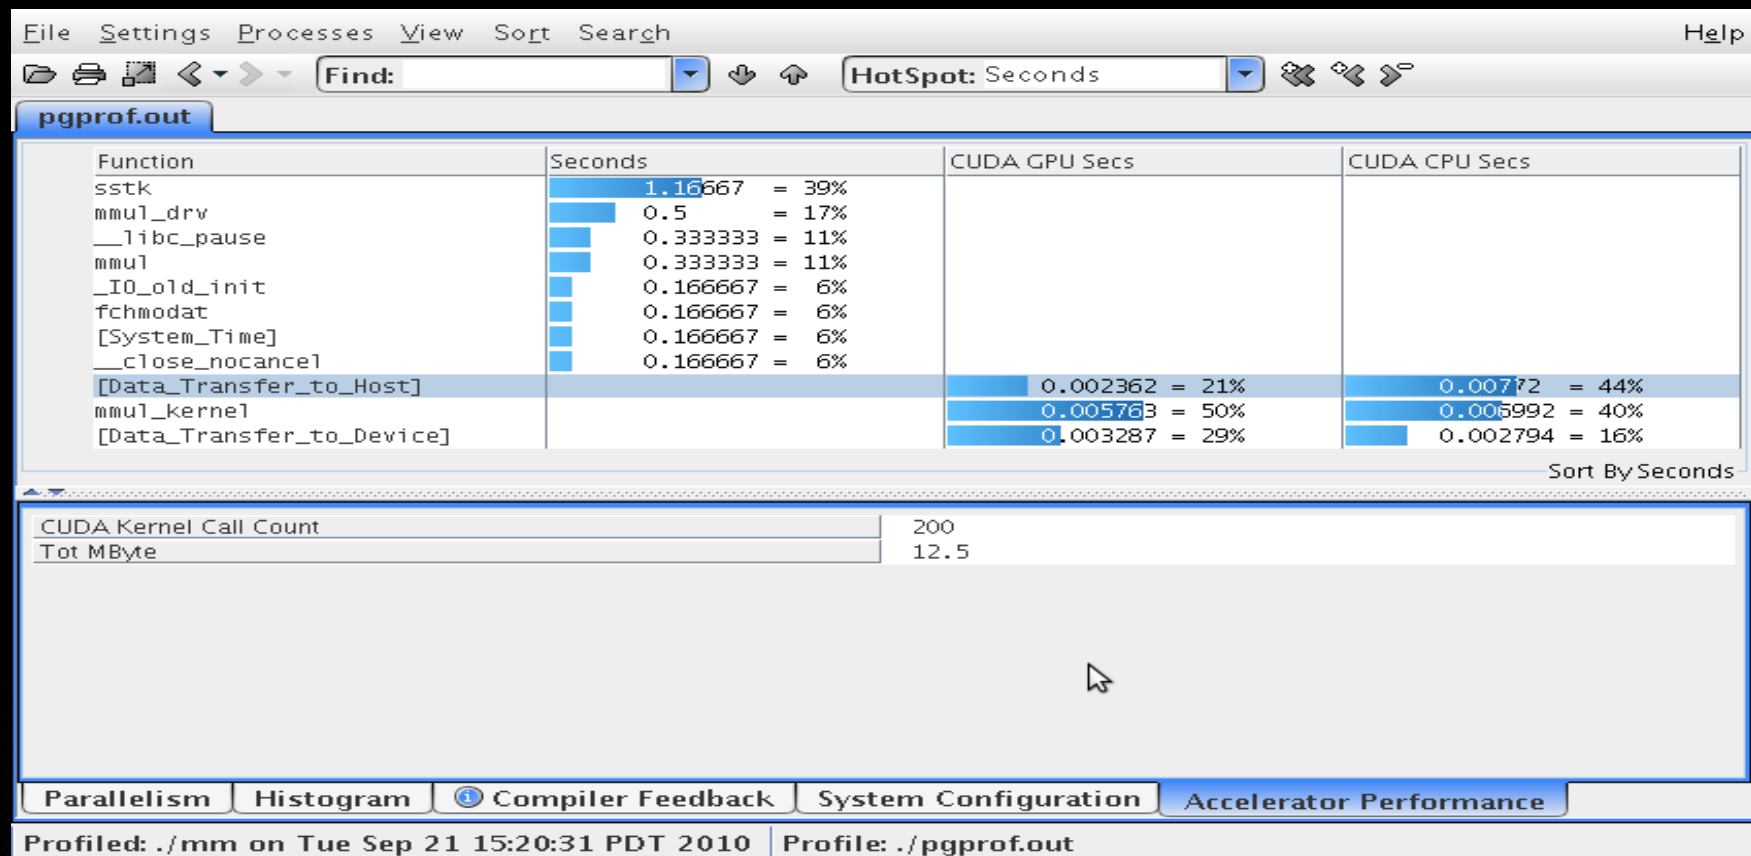
```
! timing experiment
time = 0.0
istat = cudaEventRecord(start, 0)
do j = 1, NREPS
    call sgemmNN_16x16<<<blocks, threads>>>(dA, dB, dC, m, N, k, alpha, beta)
end do
istat = cudaEventRecord(stop, 0)
istat = cudaDeviceSynchronize()
istat = cudaEventElapsedTime(time, start, stop)
time = time / (NREPS*1.0e3)
```

pgcollect / pgprof

- pgcollect [-cuda] a.out
 - Use -cuda to collect compute profile data
- pgprof [-exe a.out] [pgprof.out]







nvprof a.out

```
==30736== NVPROF is profiling process 30736, command: a.out
```

```
==30736== Profiling application: a.out
```

```
==30736== Profiling result:
```

Time (%)	Time	Calls	Avg	Min	Max	Name
38.89%	393.41us	4	98.352us	98.176us	98.496us	[CUDA memcpy HtoD]
38.88%	393.25us	4	98.312us	97.760us	98.688us	[CUDA memcpy DtoH]
6.05%	61.184us	1	61.184us	61.184us	61.184us	testasync_31_gpu
5.92%	59.872us	1	59.872us	59.872us	59.872us	testasync_41_gpu
5.88%	59.520us	1	59.520us	59.520us	59.520us	testasync_51_gpu
4.38%	44.320us	1	44.320us	44.320us	44.320us	testasync_21_gpu

nvprof -print-gpu-trace -csv a.out

Start ms	Duration us	Grid X	Grid Y	Grid Z	Block X	Block Y	Block Z	Registers Per Thread	Static SMem B	Dynamic SMem B	Size MB	Throughput GB/s	Device	Context	Stream	Name
249.4308	98.304										1	10.17253	Tesla K40c (0)	1	8	[CUDA memcpy HtoD]
250.1725	44.064	977	1	1	256	1	1	19	0	0			Tesla K40c (0)	1	8	testasync_21_gpu [22]
250.22	98.272										1	10.17584	Tesla K40c (0)	1	8	[CUDA memcpy DtoH]
269.0212	98.528										1	10.1494	Tesla K40c (0)	1	9	[CUDA memcpy HtoD]
269.6442	60.032	977	1	1	256	1	1	19	0	0			Tesla K40c (0)	1	9	testasync_31_gpu [36]
269.7065	98.272										1	10.17584	Tesla K40c (0)	1	9	[CUDA memcpy DtoH]
288.3552	98.304										1	10.17253	Tesla K40c (0)	1	10	[CUDA memcpy HtoD]

Performance Optimization

- Profile the CPU application
 - ensure you are accelerating the right part of your code
- Look at data movement
 - look for hidden data transfers
- Look at expensive kernels

Analyzing Kernels

- Launch configuration
 - block size > 1
 - grid size > 1
- Occupancy
 - registers per thread
- Memory operations
 - too many memory loads/stores (no caching)
 - stride-1 in vector index

Performance Tuning

- Performance Measurement
- Choose an appropriately parallel algorithm
- Optimize data movement between host and GPU
 - frequency, volume, regularity
- Optimize device memory accesses
 - strides, alignment
 - use shared memory, avoid bank conflicts
 - use constant memory
- Optimize kernel code
 - redundant code elimination
 - loop unrolling
 - Optimize compute intensity
 - unroll the parallel loop

Host-GPU Data Movement

- Avoid movement altogether
- Move outside of loops
- Better to move a whole array than subarray
- Update halo regions rather than whole array
 - use GPU to move halo region to contiguous area?
- Use streams, overlap data / compute
 - requires pinned host memory

Occupancy

- How many simultaneously active warps / maximum (maximum is 24, 32 (Tesla-10), 48 (Fermi), 64 (Kepler))
- Limits
 - threads per multiprocessor
 - threads per thread block - 512 (Tesla) or 1024 (Fermi/Kepler)
 - thread blocks per multiprocessor - 8 (Tesla/Fermi) or 16 (Kepler)
 - register usage - 8K / 16K / 32K / 64K registers per multiprocessor
 - each warp uses $32n$, so $16K_{reg} = 512w_{reg}$
 - shared memory usage - 16KB (Tesla) or 48KB (Fermi, Kepler)
- Low occupancy often leads to low performance
- High occupancy does not guarantee high performance

Execution Configuration

- Execution configuration affects occupancy
- Want many threads per thread block
 - multiple of 32
 - 64, 128, 192, 256
- Want many many thread blocks

Divergence

- Scalar threads executing in SIMD mode

```
if( threadIdx%x <= 10 )then
    foo = foo * 2
else
    foo = 0
endif
```

- Each path taken

```
do i = 1, threadIdx%x
    a(threadIdx%x,i) = 0
enddo
```

- Only matters within a warp

Divergence

- Pad arrays to multiples of block size

```
i = (blockidx%x-1)*64 + threadidx%x  
if( i <= N ) A(i) = ...
```

Global Memory

- Stride-1 accesses
 - address is aligned to $\text{mod}(\text{threadidx}\%x, 16)$
 - $\text{threadidx}\%x$ and $\text{threadidx}\%x+1$ access consecutive addresses
- Using shared memory as data cache
 - Redundant data access within a thread
 - Redundant data access across threads
 - Stride-1 data access within a thread

Redundant Access Within a GPU Thread

```
! threadIdx%x from 1:64
! this thread block does 256 'i' iterations
ilo = (blockidx%x-1)*256
ihi = blockidx*256 - 1
...
do j = jlo, jhi
  do i = ilo+threadidx%x, ihi, 64
    A(i,j) = A(i,j) * B(i)
  enddo
enddo
```

Redundant Access Within a GPU Thread

```
real,shared :: BB(256)
...
do ii = 0, 255, 64
  BB(threadidx%x+ii) = B(ilo+ii)
enddo
call syncthreads()
do j = jlo, jhi
  do i = ilo+threadidx%x, ihi, 64
    A(i,j) = A(i,j) * BB(i-ilo)
  enddo
enddo
```


Redundant Access Across GPU Threads

```
! threadidx%x from 1:64
i = (blockidx%x-1)*64 + threadidx%x
...
do j = jlo, jhi
  A(i,j) = A(i,j) * B(j)
enddo
```

Redundant Access Across GPU Threads

```
real, shared :: BB(64)

i = (blockidx%x-1)*64 + threadidx%x
...
do j = jlo, jhi, 64
  BB(threadidx%x) = B(jb+threadidx%x)
  call syncthreads()
  do j = jlo, min(jhi, jlo+63)
    A(i, j) = A(i, j) * BB(j-jlo+1)
  enddo
  call syncthreads()
enddo
```

stride-1 access!

Shared Memory

- 16 memory banks
- Use `threadidx%x` in leading (stride-1) dimension
- Avoid stride of 16

Unroll the Parallel Loop

- If thread 'j' and 'j+1' share data, where
 - j is a parallel index
 - j is not the stride-1 index
- Unroll two or more iterations of 'j' into the kernel

CUDA Fortran 2012/2013 Features

- PGI 2012 (CUDA 4.1 and 4.2)
 - Texture memory support
 - Support for double precision atomic add
- PGI 2013 (CUDA 5.0 and 5.5)
 - Support for dynamic parallelism
 - Relocatable device code, linking, device libraries
 - Full atomic datatype support
 - Kepler shuffle support

Latest Features in CUDA Fortran

- PGI 2014 (CUDA 6.0 and beyond)
 - DWARF generation, GPU-side debugging
 - Full shuffle datatype support
 - Overloaded reduction intrinsics
 - Support for Unified Memory
 - Improved interoperability with OpenACC

Debugging CUDA Fortran with Allinea DDT

The screenshot displays the Allinea DDT 4.1.1-33360 [Trial Version] interface. The main window shows the source code for a CUDA Fortran program, `mtrans.cuf`. The code defines a subroutine `transposeNoBankConflicts` that takes `odata` and `idata` as arguments. It uses `implicit none` and declares `real` and `integer` variables. The subroutine calculates `x` and `y` indices based on `blockIdx` and `threadIdx`. It then performs a loop over `j` from 0 to `TILE_DIM-1`, updating `idata` and `odata` using `tile` and `threadIdx` indices. The program also includes a `program transposes` block that uses `cudafor` and `use kernels_m` to execute the subroutine.

The interface includes a Project Files pane on the left, a Locals pane on the right, and an Input/Output pane at the bottom. The Input/Output pane shows the following output:

```
Process 0:
Process 0: Device: Tesla K20c
Process 0: Matrix size: 1024 1024, Block size: 32 8, Tile size: 32 32
Process 0: dimGrid: 32 32 1, dimBlock: 32 8 1
Process 0:
Process 0: Routine Bandwidth (GB/s)
```

The Locals pane shows the following variables and their values:

Variable Name	Value
<code>blockIdx</code>	<code>(x = 1, y = 1, z = 1)</code>
<code>idata</code>	<code><error reading variable idata></code>
<code>j</code>	<code><value optimized out></code>
<code>odata</code>	<code>(x = 1, y = 1, z = 1)</code>
<code>threadIdx</code>	<code>(x = 1, y = 1, z = 1)</code>
<code>tile</code>	<code>1</code>
<code>x</code>	<code><value optimized out></code>
<code>y</code>	<code><value optimized out></code>

The Input/Output pane also includes a section for "Type here ('Enter' to send):" with a "More" button.

Dwarf Generation Enables Tools

```
% pgf90 -g -O0 -Mcuda=cc35 saxpy.cuf
```

```
% cuda-memcheck ./a.out
```

```
===== CUDA-MEMCHECK
```

```
0: copyout Memcpy (host=0x68e9e0, dev=0x500200200, size=124)
```

```
FAILED: 4(unspecified launch failure)
```

```
===== Invalid __global__ read of size 4
```

```
===== at 0x00000530 in /home/brentl/saxpy.cuf:10:m_saxpy_
```

```
===== by thread (31,0,0) in block (0,0,0)
```

```
9  ! if (i<=n) y(i) = alpha * x(i) + y(i)
10  y(i) = alpha * x(i) + y(i)
```


Extending F90 Intrinsics to Device Data

- Evaluate `sum()` of device array and return result to host

```
! host code  
sum_h = sum(a_d)
```

- Typical reduction uses two kernels
 - Calculate partial sum for threads within a block
 - Launch a one-block kernel to perform a final sum

Extending F90 Intrinsics to Device Data

- **Small arrays transfer data to host and sum there**
 - Limited parallelism for small array
 - Device to host transfer is mostly overhead
- **Large arrays perform partial sum on device, final sum on host**
- **Cutoff for optimal performance depends on host system**
 - Configurable via initialization routine

Stages of Partial Sum

- Reduce to grid of threads using grid-stride loop
- Reduce within warp using SHFL (shuffle) functions
- One thread in each warp writes the value to shared memory, all call `syncthreads()`, then warp 1 reads the values into registers and uses a set of SHFL operations again
- This results in `gridDim%x` partial sums which are moved to the host

First Stage of Partial Sum

- Reduce to grid of threads using grid-stride loop

```
nGrid = gridDim%x*blockDim%x
s = 0.0
ig = (blockIdx%x-1)*blockDim%x + threadIdx%x
do i = ig, n, nGrid
    s = s + a(i)
end do
```


SHFL (shuffle) functions

- Intra-warp data exchange
- Threads can read other threads' registers
- No shared memory required, no synchronization barriers
- Requires compute capability ≥ 3.0

`s = threadIdx%x`

`s:`

1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
---	---	---	---	---	---	---	---	---	----	----	----	----	----	----	----

 ...

`t = __shfl_xor(s, 1)`

`t:`

2	1	4	3	6	5	8	7	10	9	12	11	14	13	16	15
---	---	---	---	---	---	---	---	----	---	----	----	----	----	----	----

 ...

`s = s +`
`t`

`s:`

3	3	7	7	11	11	15	15	19	19	23	23	27	27	31	31
---	---	---	---	----	----	----	----	----	----	----	----	----	----	----	----

 ...

`t = __shfl_xor(s, 2)`

`s = s + t`

`s:`

10	10	10	10	26	26	26	26	42	42	42	42	58	58	58	58
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

 ...

Reduce within warp using SHFL functions


```
! reduce within warp
t = __shfl_xor(s, 1)
  s = s + t
t = __shfl_xor(s, 2)
  s = s + t
t = __shfl_xor(s, 4)
  s = s + t
t = __shfl_xor(s, 8)
  s = s + t
t = __shfl_xor(s, 16)
  s = s + t
```

Last Stage of Partial Sum

- Reduce to $\text{gridDim}\%x$ partial sums, one from each thread block

```
if (BlockDim%x == 32) then
  if (threadIdx%x == 1) p(blockIdx%x) = s
else
  warpID = (threadIdx%x-1)/32+1
  laneID = iand(threadIdx%x,31)
  if (laneID == 1) p_s(warpID) = s
  call syncthreads()
  if (warpID == 1) then    ! Reduce in warp 1
    width = blockDim%x/32
    s = p_s(threadIdx%x)
    i = 1
    do while (i < width)
      t = __shfl_xor(s, i, width)
      s = s + t
      i = i*2
    end do
    if (threadIdx%x == 1) p(blockIdx%x) = s
  end if
end if
```


Summation Performance in Seconds



CUDA Runtime API Routines

- Device management
- Thread management
- Memory management
- Event management
- Fortran: use cudafor

Device Management

- `integer cudaGetDeviceCount(icount)`
`integer, intent(out) :: icount`
- `integer cudaSetDevice(inum)`
`integer, intent(in) :: inum`
- `integer cudaGetDevice(inum)`
`integer, intent(out) inum`

Device Management

- `integer cudaGetDeviceProperties(prop, inum)`
 `type(cudaDeviceProp), intent(out) :: prop`
 `integer, intent(in) :: inum`
- `integer cudaChooseDevice(inum, prop)`
 `type(cudaDeviceProp), intent(in) :: prop`
 `integer, intent(out) :: inum`

Device Management

- `prop%name` [character]
- `prop%major`
- `prop%minor`
- `prop%totalGlobalMem`
- `prop%regsPerBlock`
- `prop%maxThreadsPerBlock`
- `prop%maxThreadDim(3)`
- `prop%maxGridSize(3)`
- `prop%clockRate`
- `prop%totalConstMem`

Device Management

- `prop%warpSize`
- `prop%textureAlignment`
- `prop%deviceOverlap`
- `prop%multiProcessorCount`
- `prop%kernelExecTimeoutEnabled`
- `prop%memPitch`
- `prop%integrated`
- `prop%canMapHostMemory`
- `prop%computeMode`

Stream Management

- `integer cudaStreamCreate(stream)`
`integer, intent(out) :: stream`
- `integer cudaStreamQuery(stream)`
`integer, intent(in) :: stream`
- `integer cudaStreamSynchronize(stream)`
`integer, intent(in) :: stream`
- `integer cudaStreamDestroy(stream)`
`integer, intent(in) :: stream`

Memory Management

- `integer cudaMallocHost(ptr, size)`
 `type(c_ptr),intent(out) :: ptr`
 `integer,intent(in) :: size`
- `integer cudaFreeHost(ptr)`
 `type(c_ptr),intent(in) :: ptr`
- `integer cudaMalloc(ptr, count)`
 `<type>,device,allocatable,dimension(:) ::&`
 `ptr`
 `integer,intent(in) :: count`
- `integer cudaFree(ptr)`
 `<type>,device,dimension(*) :: ptr`

Memory Management

- `integer cudaMemcpy(dst,src,count,dir)`
`<type>,<device><,dimension> :: dst,src`
`integer, intent(in) :: count, dir`
- **dir values:**
`cudaMemcpyHostToHost`
`cudaMemcpyHostToDevice`
`cudaMemcpyDeviceToHost`
`cudaMemcpyDeviceToDevice`

Memory Management

- `integer cudaMallocPitch(ptr, pitch, w, h)`
`<type>,device,allocatable,dimension(:,:)&`
`:: ptr`
`integer,intent(out) :: pitch`
`integer,intent(in) :: w,h`
- `integer cudaMemset(ptr, value, count)`
`<type>,device<,dimension(*)> :: ptr`
`<type> :: value`
`integer :: count`
- `integer cudaMemset2D(ptr,pitch,value,w,h)`
`<type>,device<,dimension> :: ptr`

CUDA Fortran with Managed Data

```
program pgi14x
  use cudafor
  integer, parameter :: n = 100000
  integer, managed :: a(n), b(n), c(n)
  a = [(i,i=1,n)]
  b = 1
  !$cuf kernel do <<< *, * >>>
    do i = 1, n
      c(i) = a(i) + b(i)
    end do
    if (sum(c).ne.n*(n+1)/2+n) then
      print *,c(1),c(n)
    end if
  end
```

What is CUDA Managed Data?

```
real, managed, allocatable, dimension(:,:) ::  
    A,B,C  
  
...  
  
allocate (A(N,M), B(M,L), C(N,L))  
...  
  
call mm_kernel <<<dim3(N/16,M/16),dim3(16,16)>>>  
    ( A, B, C, N, M, L)  
  
deallocate ( A, B, C )  
  
...
```

Host Code

```
attributes(global) subroutine mm_kernel  
    ( A, B, C, N, M, L )  
real :: A(N,M), B(M,L), C(N,L), Cij  
integer, value :: N, M, L  
integer :: i, j, kb, k, tx, ty  
real, shared :: Asub(16,16), Bsub(16,16)  
tx = threadidx%x  
ty = threadidx%y  
i = blockidx%x * 16 + tx  
j = blockidx%y * 16 + ty  
Cij = 0.0  
do kb = 1, M, 16  
    Asub(tx,ty) = A(i,kb+tx-1)  
    Bsub(tx,ty) = B(kb+ty-1,j)  
    call syncthreads()  
    do k = 1,16  
        Cij = Cij + Asub(tx,k) * Bsub(k,ty)  
    enddo  
    call syncthreads()  
enddo  
C(i,j) = Cij  
end subroutine mmul_kernel
```

Device Code

Managed Data enables Derived Type Usage

```
type tCM
  integer, allocatable, device :: fine(:)
  real, allocatable, device :: mat_matrix(:, :, :)
  real, allocatable, device :: src_matrix(:, :, :)
end type
type(tCM), allocatable, managed :: cm_list(:)
```

Array of derived type can be allocated on the host and used on the device

CUDA Fortran/OpenACC Interoperability

- Calling CUDA Fortran kernels from accelerator data regions
- Using CUDA Fortran device data in compute regions

```
!$acc data create(x) copy(y)
a = 3.0
!$acc kernels loop
do i = 1, n
    x(i) = x_dev(i) + 1.0
    y(i) = y(i) + real(i)
end do
call mysaxpy <<<block, thread>>> (n, a, x, y)
!$acc end data
```

- Support of arguments with the device attribute in our OpenACC 2.0 Runtime Library Routines

Calling CUDA Fortran Subroutines From OpenACC Programs

```
use cublas
```

```
!$acc data region copyin(x)
```

```
! some compute regions . . .
```


```
k = isamax(N,x,1)
```

```
! maybe some other compute regions. . .
```

```
!$acc end data region
```

- You can call CUDA routines by creating explicit interfaces to host-resident data and device-resident data specific functions for a generic call

CUDA Fortran Literature



- Book Released in 2013
- PGInsider articles at <http://www.pgroup.com>
- CUDA Fortran Reference Manual